

Výkonové optimalizace MySQL

Anna Janáčková, Petr Medonos

O čem to bude?

```
./mysqld --run-fast
```

O čem to bude?



O čem to bude?

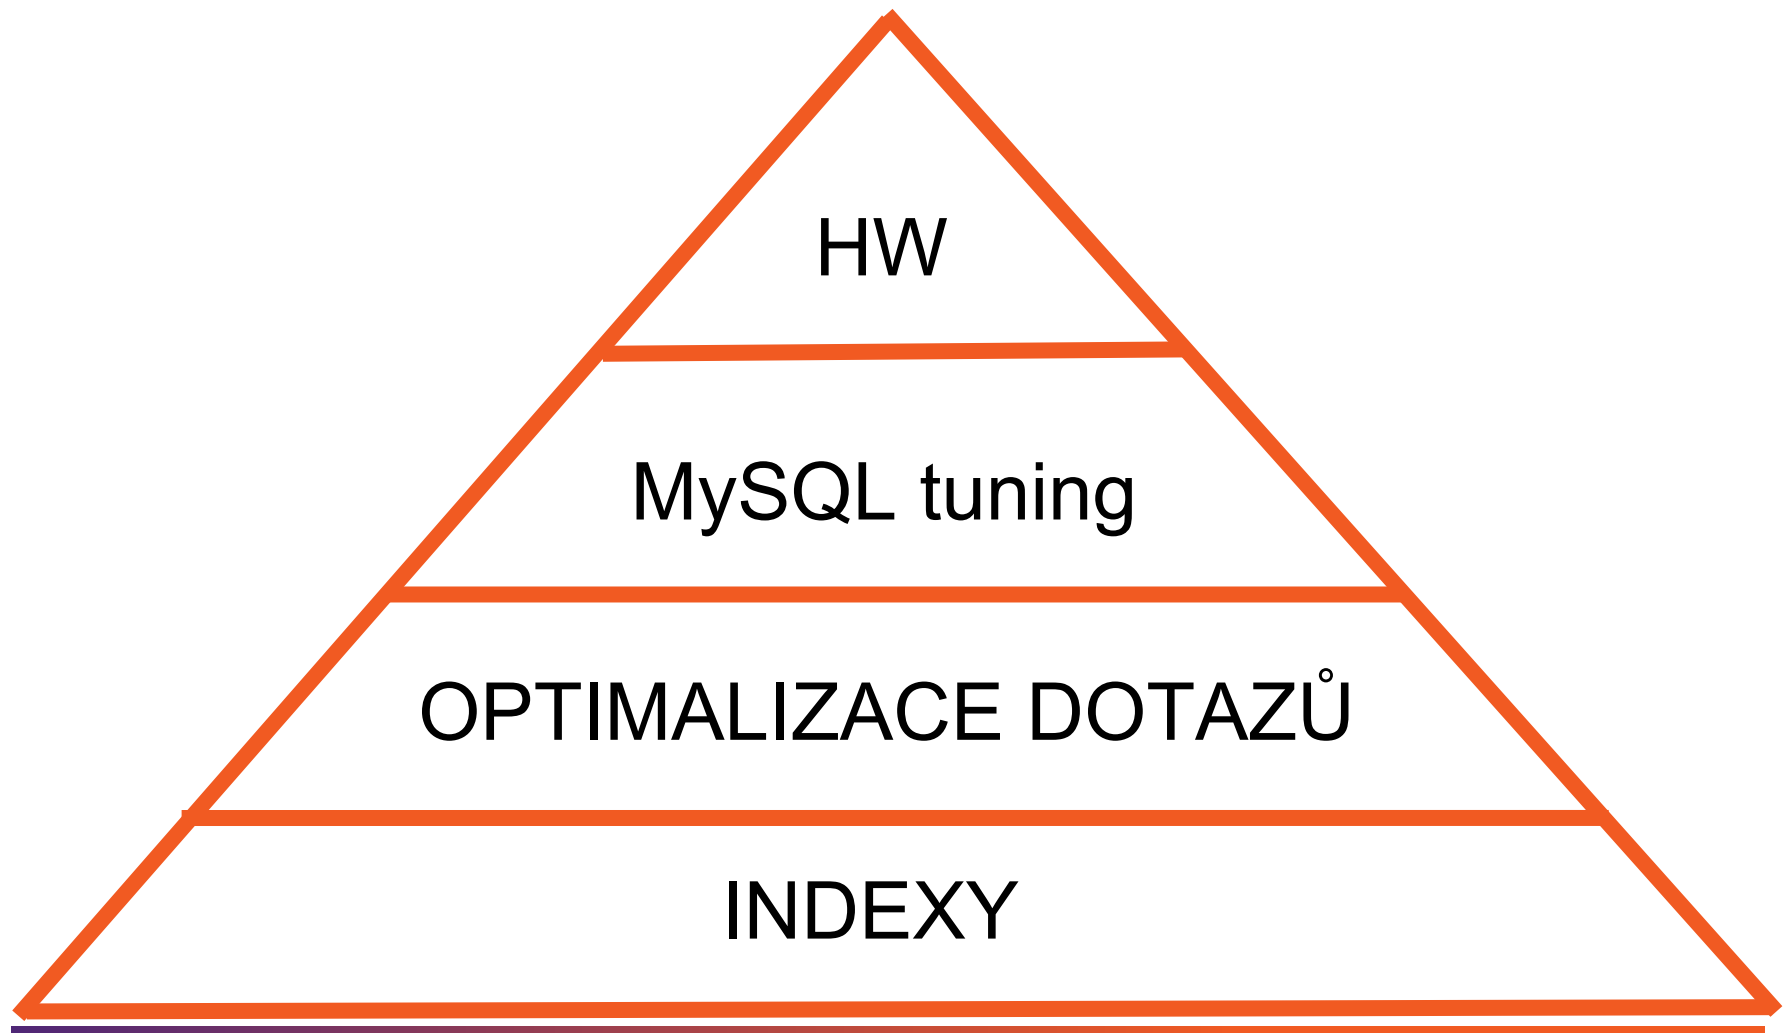
- výběr HW
- tuning GNU/Linux
- optimalizace paramerů MySQL
- analýza provozu a řešení problémů

MySQL v ETN



- 4.1, 5.0, 5.1, 5.5
- datasety ~100MB - ~10GB
- 100% Linux
- x86_64, Intel

Efektivita ladění



Paretovo pravidlo

80% - 20%

Čísla, které by měl každý znát ...

<http://odbms.org/download/dean-keynote-ladis2009.pdf>

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	25 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	3,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from disk	20,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns

Hardware



Hardware

● CPU

- MySQL neumí rozkládat dotazy na víc jader
 - => výkon CPU je důležitý
- každé spojení má vlastní thread
 - => dostatečný počet jader

● RAM

- OLTP - aktivně používaná data se musí vejít do RAM

● SÍŤ

- komunikace s aplikací
- replikace
 - dedikovaná síť, bonding 2x 1Gbit


Hardware - I/O

- potřebujeme maximální rychlost a **minimální latenci**
 - **raid10**
 - hardwarový řadič
 - “zadrátovaná” logika je rychlejší
 - baterie pro **writeback cache**
 - vypnout write barriers
`mount -o nobarrier`

GNU/Linux

- read-ahead
 - `blockdev --setra <#> /dev/<...>`
- vhodný filesystem
 - XFS
 - ext4
- I/O scheduler - deadline
- snížit swappiness
 - `echo "20" > /proc/sys/vm/swappiness`

Optimalizace parametrů MySQL serveru



Alokace paměti

```
set @giga=1024*1024*1024;  
SELECT ( @@key_buffer_size +  
  @@query_cache_size +  
  @@innodb_buffer_pool_size +  
  @@innodb_additional_mem_pool_size +  
  @@innodb_log_buffer_size + @@max_connections  
  * (  
    least(@@tmp_table_size,@@max_heap_table_size  
    ) + @@read_buffer_size +  
    @@read_rnd_buffer_size + @@sort_buffer_size  
    + @@join_buffer_size + @@binlog_cache_size +  
    @@thread_stack) ) / @giga AS MAX_MEM;
```

Fiktivní MySQL instance

- 24 core
 - 64GB RAM
 - 6x SAS 15k, raid 10 512MB cache, write-back
 - ~ 10 schémat
 - ~ 1000 tabulek
 - InnoDB only
 - dataset 25GB
 - 80% čtení - 20% zápis
-

Parametry MySQL serveru

`max_connections = 400`

`back_log = 1024`

`thread_cache_size = 300`

`skip_name_resolve = ON`


`tmp_table_size = 128MB`

`max_heap_table_size = 128MB`

`sort_buffer_size = 8MB`

`join_buffer_size = 8MB`

Query cache



- co nelze cachovat
 - `now()`, `rand()`, `uuid()`
 - tabulka s právy na úrovni sloupců
 - ...

`query_cache_size = 64-256MB`

`query_cache_limit = 16MB`

`query_cache_min_res_unit =`

Query cache

$$\text{QC_min_res} = \frac{\text{query_cache_size} - \text{Qcache_free_memory}}{\text{Qcache_queries_in_cache}}$$

```
show global variables like 'query_cache_size';  
show global status like 'Qcache_free_memory';  
show global status like 'Qcache_queries_in_cache';
```

Parametry InnoDB engineu

<code>innodb_buffer_pool_size</code>	<code>= 32G</code>
<code>innodb_buffer_pool_instances</code>	<code>= 32</code>
<code>innodb_purge_threads</code>	<code>= 1</code>
<code>innodb_log_file_size</code>	<code>= 2G</code>
<code>innodb_flush_method</code>	<code>= O_DIRECT</code>
<code>innodb_io_capacity</code>	<code>= 10000</code>
<code>innodb_io_capacity_max</code>	<code>= 15000</code>
<code>innodb_lru_scan_depth</code>	<code>= 10000</code>

Analýza provozu a řešení problémů



Náročné operace pro MySQL

- query cache miss
- table scan, full join
- tmp tabulka
- tmp tabulka vytvořená na disku, filesort na disku

Slow query

```
slow-query-log-file      =  
    /var/log/mysql/mysql-slow.log
```

```
slow-query-log           = 1
```

```
slow_launch_time        = 1000
```

```
log-queries-not-using-indexes
```

```
long_slow_filter =
```

```
"tmp_table_on_disk, filesort_on_disk, full_joi  
n"
```

Slow query - pt-query-digest

ukázka

Detekce zámků (na úrovni InnoDB)

```
SELECT r.trx_id waiting_trx_id, r.trx_mysql_thread_id  
       waiting_thread, p1.USER waiting_user, p1.HOST  
       waiting_host, r.trx_query waiting_query, b.trx_id  
       blocking_trx_id, b.trx_mysql_thread_id  
       blocking_thread, p2.USER blocking_user, p2.HOST  
       blocking_host, b.trx_query blocking_query FROM  
information_schema.innodb_lock_waits w INNER JOIN  
information_schema.innodb_trx b ON b.trx_id =  
w.blocking_trx_id INNER JOIN  
information_schema.innodb_trx r ON r.trx_id =  
w.requesting_trx_id INNER JOIN  
information_schema.PROCESSLIST p1 ON p1.id =  
r.trx_mysql_thread_id INNER JOIN  
information_schema.PROCESSLIST p2 ON p2.id =  
b.trx_mysql_thread_id;
```


Nástroje optimalizace

- explain
- profiling
- performance schema
- (show status, show engine innodb status)

Explain

- jak je dotaz vyhodnocen?
- jaké indexy se používají?

```
explain select SQL_NO_CACHE item.id  
  from item JOIN item_review on  
    item.id=item_review.item_id group  
  by item.id order by item.id\G;
```

Explain

```
        id: 1
select_type: SIMPLE
        table: item_review
      type: ALL
possible_keys: NULL
        key: NULL
      key_len: NULL
        ref: NULL
        rows: 41310
      Extra: Using temporary; Using
filesort; Start temporary
```

Explain

```
        id: 1
select_type: SIMPLE
        table: item_review
        type: index
possible_keys: item_order
        key: item_order
    key_len: 8
        ref: NULL
        rows: 41310
        Extra: Using index; Using temporary;
Using filesort
```

Profilování

```
select SQL_NO_CACHE distinct id from item  
where id IN (select item_id from  
item_review) order by id;
```

9038 rows in set (1.31 sec)

```
select SQL_NO_CACHE item.id from item JOIN  
item_review on item.id=item_review.item_id  
group by item.id order by item.id;
```

9038 rows in set (0.07 sec)

Profilování

```
set profiling = 1;  
select Q1 ...  
select Q2 ...  
set profiling = 0;
```

```
show profiles;  
show profile for query 1;  
show profile for query 2;
```

Profilování

Status	Duration 1	Duration 2
optimizing	0.000007	0.000006
Creating tmp table	0.000000	0.000027
executing	0.000006	0.000003
Copying to tmp table	0.000000	0.065036
Sorting result	0.000004	0.001362
Sending data	0.000038	0.002149
optimizing	0.000007	0.001362
statistics	0.000013	0.002149
preparing	1.306187	0.000005

Performance schéma

- MySQL 5.5+
- systémová databáze
- SQL rozhraní

```
performance_schema = ON  
show variables like 'perf%';  
show engine performance_schema  
status;
```


Performance schéma

- základní nastavení pomocí `setup_` tabulek
 - `setup_instruments`
 - `setup_consumers`
 - `setup_timers`
 - `setup_actors`
 - `setup_objects`
- `ps_helper` (Mark Leith)

Performance schéma - instruments

```
wait/io/file/sql/binlog  
wait/io/file/sql/binlog_index  
...  
statement/sql/select  
statement/sql/update  
statement/sql/insert  
...  
stage/sql/Opening tables  
stage/sql/optimizing  
stage/sql/preparing
```

Performance schéma (ps_helper)

```
mysql> select * from schema_unused_indexes where  
      object_schema != 'mysql' and index_name !=  
      'PRIMARY';
```

object_schema	object_name	index_name
testcz	item	title_cz
testcz	item	url_id_cz
testcz	item	item_segment_id
testcz	item_review	PRIMARY
testcz	item_review	eid

Optimalizační hinty

- SQL_NO_CACHE/SQL_CACHE
- USE/FORCE/IGNORE INDEX
- STRAIGHT_JOIN
- SQL_BIG_RESULT
-

Optimalizační hinty

```
SELECT SQL_CALC_FOUND_ROWS STRAIGHT_JOIN  
  DISTINCT item.id FROM  item JOIN  
  item_available ON
```

```
SELECT * FROM table1 USE INDEX  
  (col1_index,col2_index)  
  WHERE col1=1 AND col2=2 AND col3=3;
```

Co když to nestačí ...


- replikace
 - master - slave, master - master
- Cluster
 - Galera (Percona XtraDB Cluster), MySQL cluster
- sharding

Užitečné tooly

- percona toolkit
 - pt-query-digest
 - pt-mext
- ps_helper
- mysqltunner.pl
- mysql_primer
- munin mysql (<https://github.com/kjellm>)
- percona nagios plugins

Zdroje

- High Performance MySQL (3. edice)
- <http://www.mysqlperformanceblog.com/>
- <http://blog.jcole.us/>
- <http://dev.mysql.com/doc/refman/5.6/en/>



COME TO THE
DARK SIDE



WE HAVE COOKIES

Q&A

petr.medonos@etnetera.cz

anna.janackova@etnetera.cz